

Advanced Databases

Concurrency Control Techniques I

Dr. George Mertzios
Michaelmas Term

george.mertzios@durham.ac.uk

Room 2066, MCS Building

Tel: 42 429

Course Outline

- Enhanced Entity-Relationship (EER) Model
- Semistructured Databases - XML
- XML Data Manipulation - XPath, XQuery
- Transactions and **Concurrency Control**
- Distributed Transactions
- Distributed Concurrency Control

Schedules

- Serial schedule: a schedule where the operations of any two transactions are not interleaved
- Non-serial schedule: a schedule where the operations of some transactions are interleaved
- A non-serial schedule is (conflict) serializable if:
 - it produces a database state that can be produced by some serial execution of the same transactions
- How can we achieve serializability in practice?

Concurrency control

Two main approaches to ensure consistency when executing transactions concurrently:

- Conservative (pessimistic) methods:
 - actively avoid conflicts
 - delay (or restart) transactions when they are in conflict
- Optimistic methods:
 - assume that transactions are rarely in conflict
 - check for conflicts just before the transaction commits (i.e. at the end)

Concurrency control

- Two main **conservative** methods:
 - locking
 - timestamping
- Lock: when a transaction accesses the database, the “lock” denies access to other transactions, to prevent incorrect results

The most widely used method to *ensure* serializability

The locking method



- A transaction T can keep two types of locks:
 - shared lock (or read lock):
 - T is allowed only to **read** some data item
 - **any other** transaction can only **read** this item
 - exclusive lock (or write lock):
 - T is allowed to **read** and **write** on some data item
 - **any other** transaction has **no access** to this item
- Terminology:
 - a transaction **requests** a lock to the DBMS
 - the DBMS **grants** the lock; otherwise the transaction **waits**
 - a transaction **releases** a lock on a data item
(the item “**unlocks**”)

Rules for locks

- When a transaction needs to access a data item, it requests:
 - a **shared** lock for **read** only access
 - an **exclusive** lock for **read and write** access
- If the item is currently *not* locked by another transaction, the lock will be granted
- If the item is currently *locked*, the DBMS checks compatibility between requested lock / existing lock:
 - a **shared** lock is requested on an item already locked by a **shared** lock \Rightarrow the new lock is **granted**
 - **otherwise**: the transaction must **wait** until the existing lock is released

Rules for locks

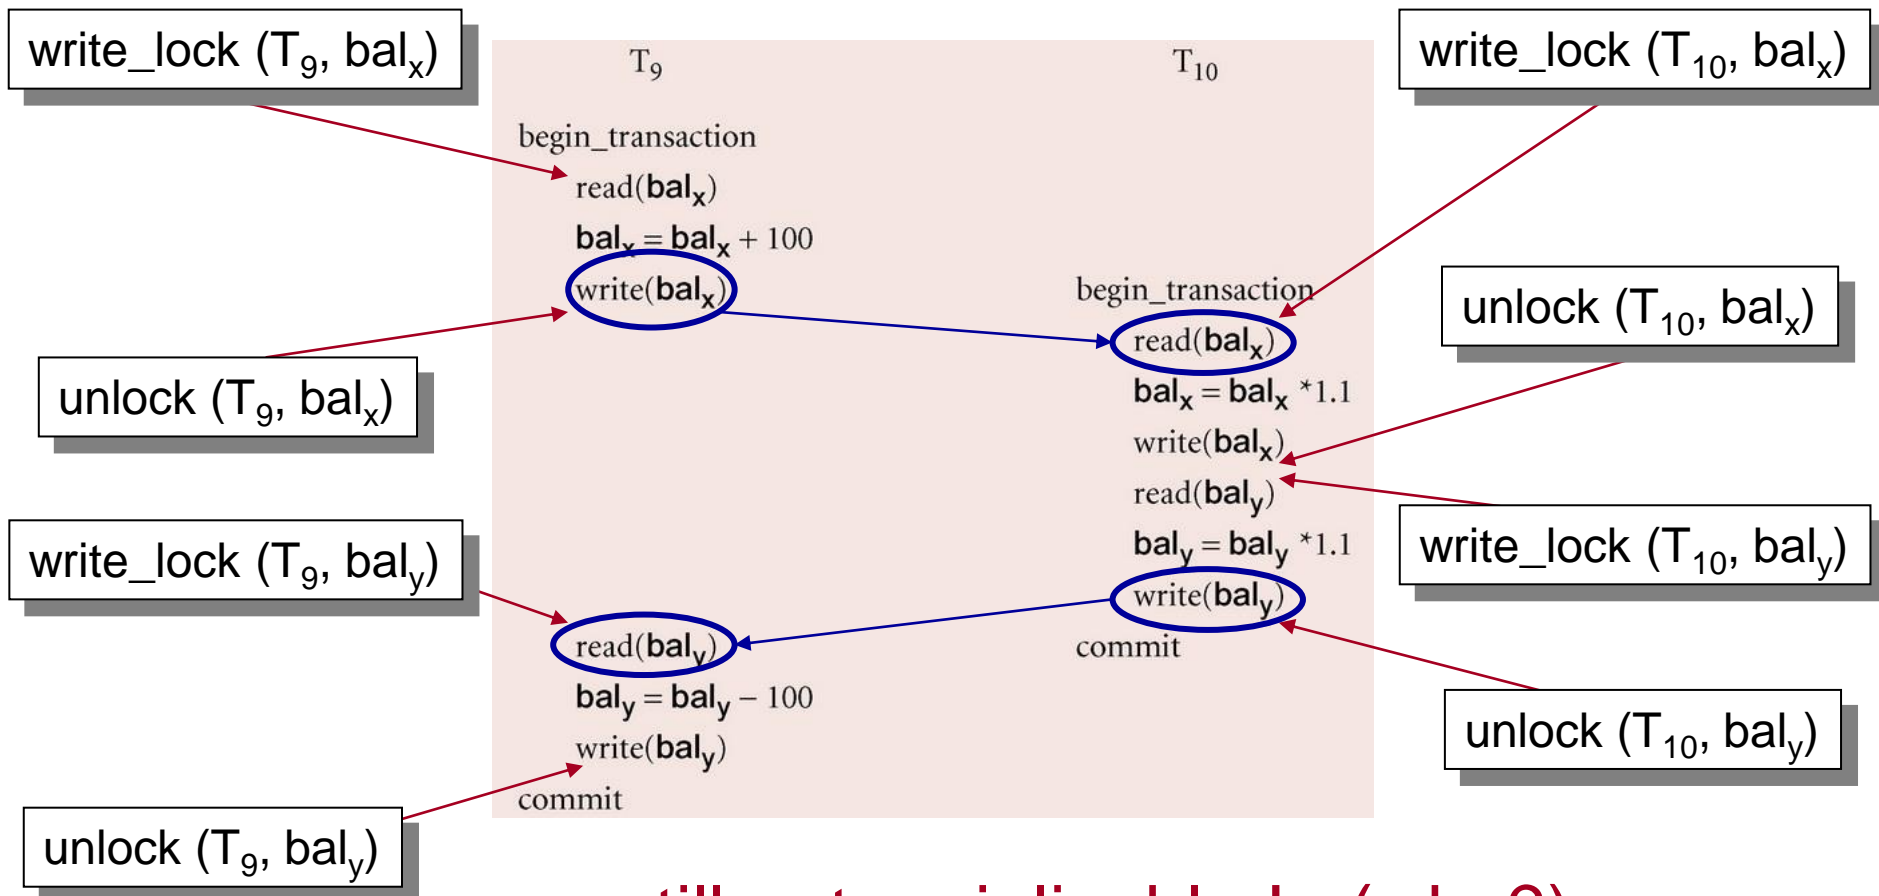
- A transaction **holds** a lock until it **explicitly releases** it:
 - during its **execution**, or
 - when it terminates, i.e.
 - when it **commits**, or
 - when it **aborts**
- The effects of a **write** operation are made **visible**:
 - only when the exclusive lock is **released**
(to ensure the **Isolation** property)
- Some systems permit a lock to be:
 - upgraded from a shared lock to an exclusive lock, or
 - downgraded from an exclusive lock to a shared lock

(to increase more the concurrency \Rightarrow efficiency)

The locking method

- Can we guarantee serializability by just using locks?

NO !



still not serializable ! (why?)

The locking method

- Can we guarantee serializability by just using locks?

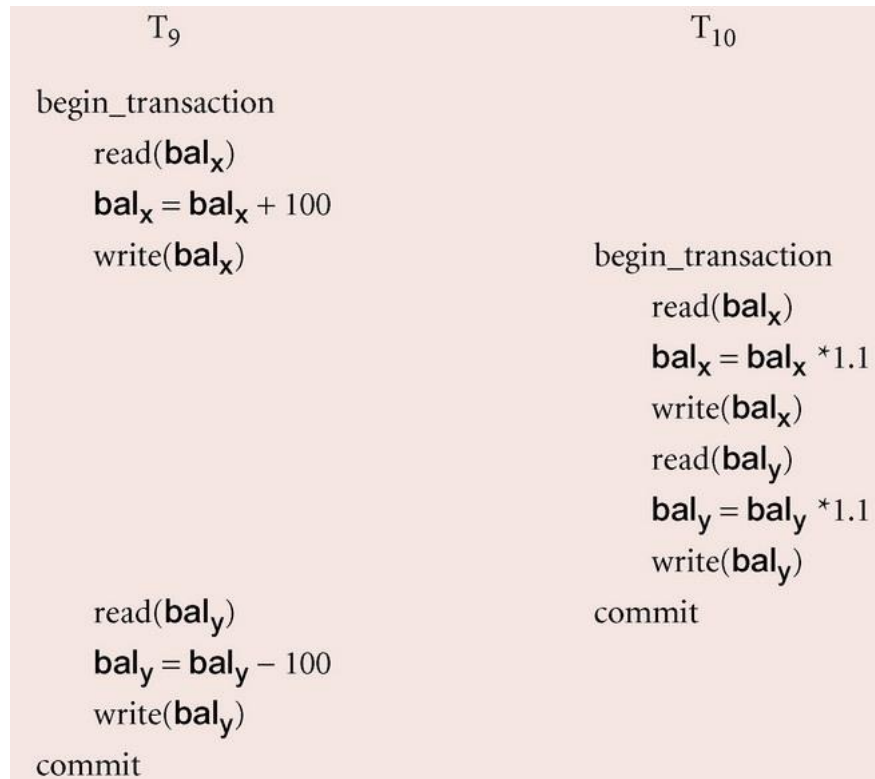
Suppose initially: $bal_x=100$, $bal_y=400$

If First T_9 , then T_{10} :

- $bal_x=220$
 $bal_y=330$
($bal_x + bal_y = 550$)

First T_{10} , then T_9 :

- $bal_x=210$
 $bal_y=340$
($bal_x + bal_y = 550$)



But in this schedule:

- $bal_x=\underline{220}$
 $bal_y=\underline{340}$
($bal_x + bal_y = \underline{560}$)

Thus:

not the same as
any serial schedule!

still not serializable ! (why?)

The locking method

- To guarantee serializability, we need an additional protocol controlling the **positioning of locks**
- **Two phase locking (2PL):**
 - for every single transaction, all **locking** operations occur **before** all **unlocking** operations
- Two phases for every transaction:
 - **growing** phase
 - acquire all needed locks / no unlock
 - **shrinking** phase
 - release the locks / no new lock

Two phase locking (2PL)

- Prevent the **lost update** problem:
 - the **request** of T_1 for an exclusive lock **waits** until the exclusive lock is **released** by T_2

Time	T_1	T_2	bal_x
t_1		begin_transaction	100
t_2	begin_transaction	write_lock(bal_x)	100
t_3	write_lock(bal_x)	read(bal_x)	100
t_4	WAIT	$bal_x = bal_x + 100$	100
t_5	WAIT	write(bal_x)	200
t_6	WAIT	commit/ unlock(bal_x)	200
t_7	→ read(bal_x)		200
t_8	$bal_x = bal_x - 10$		200
t_9	write(bal_x)		190
t_{10}	commit/ unlock(bal_x)		190

Two phase locking (2PL)

- Prevent the **dirty data** problem:
 - the **request** of T_3 for an exclusive lock **waits** until the exclusive lock is **released** by T_4
 - this happens only after the **rollback** of T_4 is completed

Time	T_3	T_4	bal_x
t_1		begin_transaction	100
t_2		write_lock(bal_x)	100
t_3		read(bal_x)	100
t_4	begin_transaction	$bal_x = bal_x + 100$	100
t_5	write_lock(bal_x)	write(bal_x)	200
t_6	WAIT	rollback/unlock(bal_x)	100
t_7	→ read(bal_x)		100
t_8	$bal_x = bal_x - 10$		100
t_9	write(bal_x)		90
t_{10}	commit/unlock(bal_x)		90

Two phase locking (2PL)

- Prevent the **inconsistent analysis** problem:

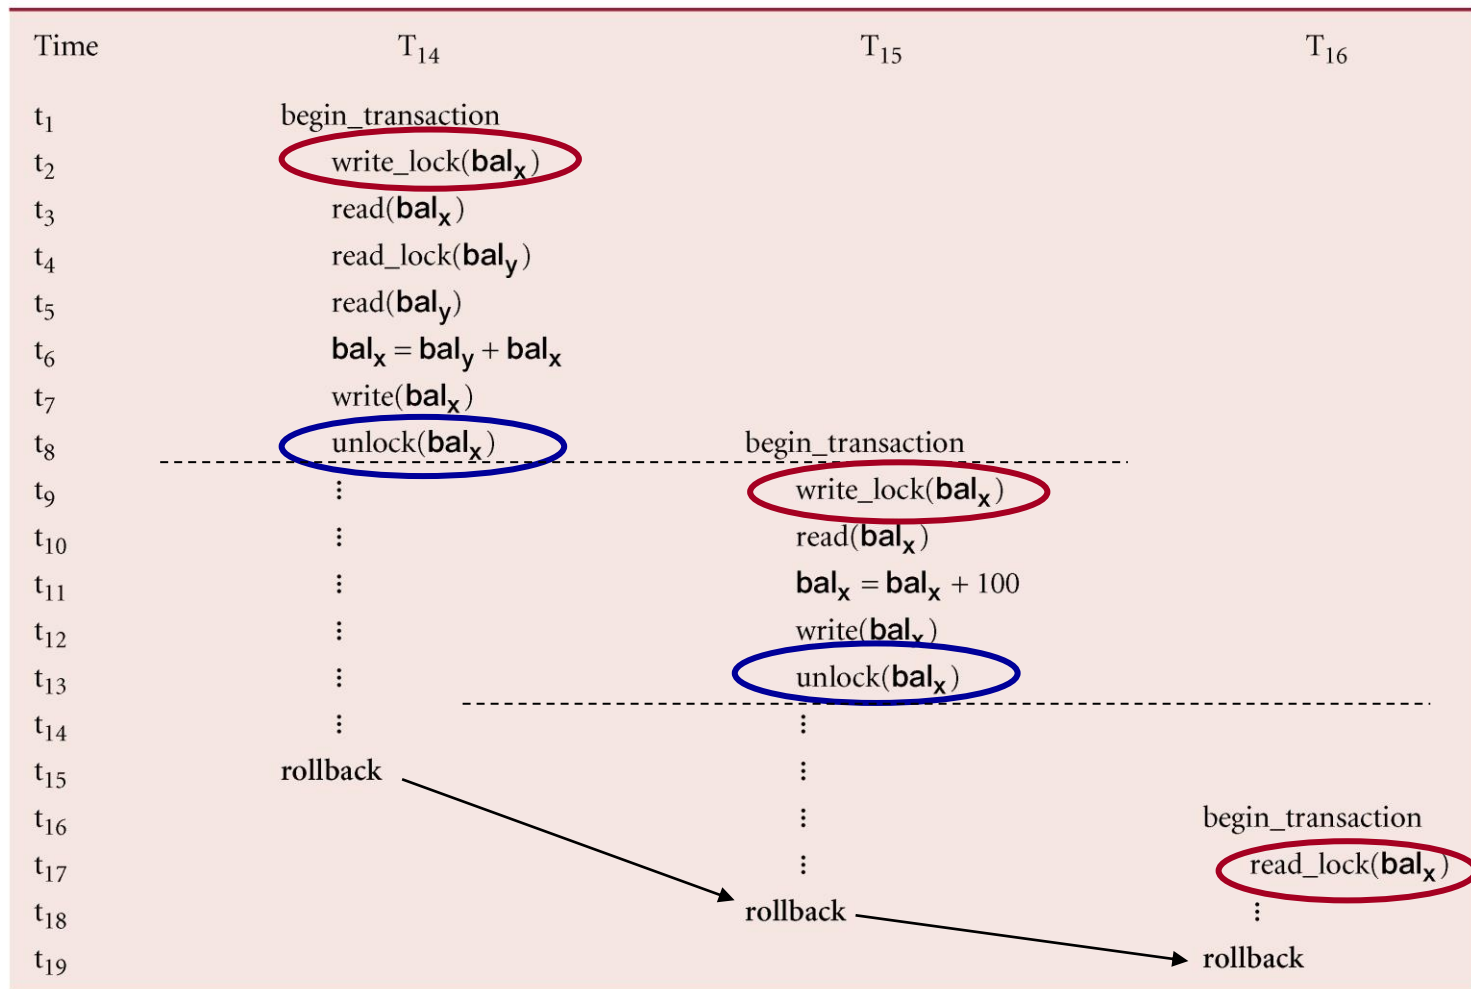
		exclusive locks		read locks					
Time	T ₅		T ₆			bal _x	bal _y	bal _z	sum
t ₁			begin_transaction			100	50	25	
t ₂	begin_transaction		sum = 0			100	50	25	0
t ₃	write_lock(bal _x)					100	50	25	0
t ₄	read(bal _x)		read_lock(bal _y)			100	50	25	0
t ₅	bal _x = bal _x - 10		WAIT			100	50	25	0
t ₆	write(bal _x)		WAIT			90	50	25	0
t ₇	write_lock(bal _z)		WAIT			90	50	25	0
t ₈	read(bal _z)		WAIT			90	50	25	0
t ₉	bal _z = bal _z + 10		WAIT			90	50	25	0
t ₁₀	write(bal _z)		WAIT			90	50	35	0
t ₁₁	commit/unlock(bal _x , bal _z)		WAIT			90	50	35	0
t ₁₂			read(bal _x)			90	50	35	0
t ₁₃			sum = sum + bal _x			90	50	35	90
t ₁₄			read_lock(bal _y)			90	50	35	90
t ₁₅			read(bal _y)			90	50	35	90
t ₁₆			sum = sum + bal _y			90	50	35	140
t ₁₇			read_lock(bal _z)			90	50	35	140
t ₁₈			read(bal _z)			90	50	35	140
t ₁₉			sum = sum + bal _z			90	50	35	175
t ₂₀			commit/unlock(bal _x , bal _y , bal _z)			90	50	35	175

Two phase locking (2PL)

- It can be formally proved that:
 - if **every transaction** in a schedule follows the **two-phase locking (2PL)** protocol, then the schedule is always **conflict serializable**
- Still:
 - **problems** can occur by **early release** of locks
- The “**cascading rollback**” problem:
 - a transaction rollbacks after long time
 - this can cause a pile up of rollbacks

⇒ **inefficient database !**

The cascading rollback problem

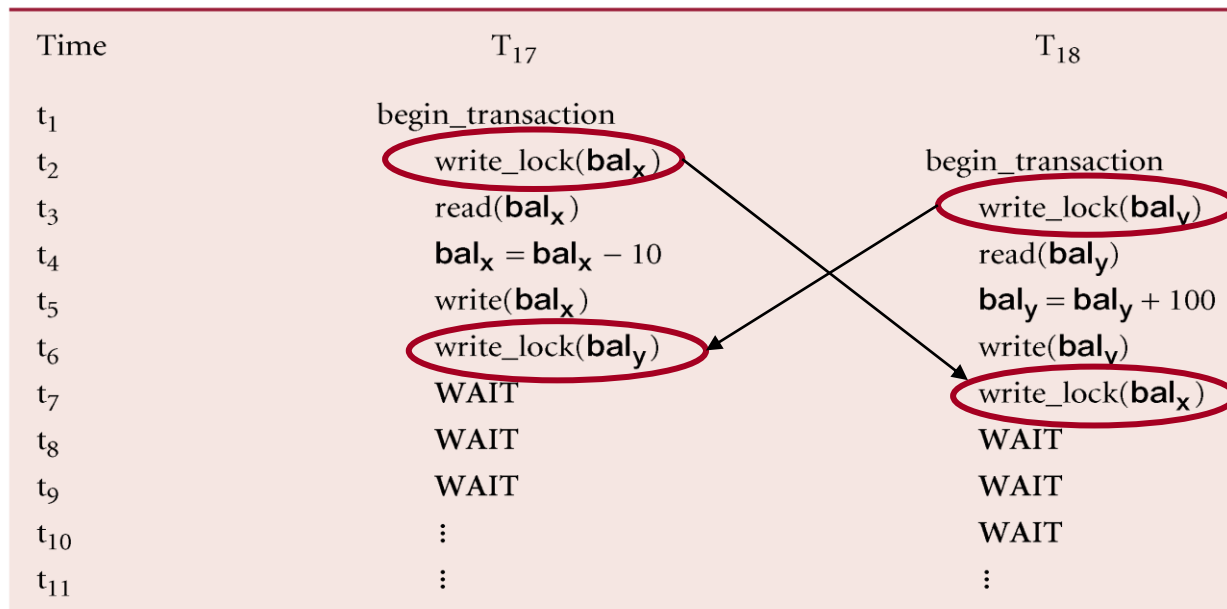


The cascading rollback problem

- Solution: the rigorous two-phase locking (2PL)
 - release all locks at the end of every transaction
(when it *commits*)
- With rigorous 2PL:
 - the transactions are serializable in the order they commit
 - no cascading rollback !
- Another variant: the strict two-phase locking (2PL)
 - release all write locks at the end of every transaction
(when it *commits*)
 - read locks can be released earlier

Deadlocks

- Deadlocks: another problem with the locking method
 - two (or more) transactions wait for each other
 - \Rightarrow they can wait for ever
 - this happens also with the 2PL protocol!
 - also when all locks are released at the end of transactions!



The DBMS has to **recognize** and **resolve** the deadlock

Handling deadlocks

Two general techniques:


- **Timeouts:**
 - a transaction requests a lock only for a system-defined (maximum) period of time
- After that point:
 - the DBMS assumes there is a deadlock
 - the request **times out**
 - the transaction **rolls back** and **restarts**
- a simple solution
- used by several commercial DBMSs



Handling deadlocks

Two general techniques:

- **Deadlock detection:**
 - construct the **wait-for graph** $G = (N, E)$
 - nodes N : one node for each transaction T_i
 - a directed edge $T_i \rightarrow T_j$ whenever transaction T_i **is waiting** for a lock that is kept by transaction T_j
- **Theorem:** a **deadlock** exists if and only if the wait-for graph contains a **directed cycle**



Handling deadlocks

- Once a **deadlock** is detected:
 - the DBMS **aborts** (**rolls back**) a transaction
 - check for a deadlock too often
 - ⇒ large computational overhead (slower database)
 - check for a deadlock too rarely
 - ⇒ deadlocks may be undetected for long periods

Important parameters for recovery:

- **How far to roll back a transaction?**
 - simplest solution: undo all changes (restart)
 - more efficient: possibly roll back only a part of the transaction

Handling deadlocks

Other important parameters for recovery:

- Choice of the deadlock “victim”:
 - which transaction to abort?
 - the choice affects efficiency of the database
- General criteria:
 - how long a transaction has been running?
 - better to abort a “short” transaction
 - how many data items did it update so far?
 - better to abort transactions that made few changes
 - how many data items does it still have to update?
 - better not to abort transactions that have few more data to update
 - difficult for the DBMS to know in advance

Handling deadlocks

Other important parameters for recovery:

- Avoiding “starvation”:
 - starvation occurs when a specific transaction is always chosen as the “victim”
 - this transaction can never complete
- Common solution:
 - store how many times every transaction has been aborted
 - when an upper limit is reached, use different selection criteria

Granularity of data items

- Further locking parameter: **granularity**
 - how “**large**” is the “**data item**” to be locked each time?
- Hierarchy of granularity:
 - **entire database** – **coarsest size** of a data item
 - **file** (relation / table)
 - **page** (section of physical disk where relations are stored)
 - **record** (tuple of relation)
 - **field value** (cell of a tuple) – **finest size** of a data item

Granularity of data items

- **Coarser data item size**
 - large locked items / fewer locks requested
 - ⇒ **lower degree** of **concurrency** permitted
- **Finer data item size**
 - small locked items / more locks requested
 - ⇒ **more locking information** needs to be stored
- Granularity affects efficiency:
 - best item size depends on the nature of transactions
(i.e. which / how many data items needed per transaction)

Concurrency control

- Next lecture:
 - **timestamping**
 - the alternative method to guarantee serializability
 - no locks \Rightarrow no deadlocks
- Locking method:
 - make conflicting transactions wait
- Timestamping method:
 - roll back and restart conflicting transactions

Summary of the Lecture

- Concurrency control methods:
 - the locking method
 - shared / exclusive locks
 - two phase locking (2PL)
 - growing / shrinking phase
 - use 2PL to prevent:
 - the lost update problem
 - the dirty data problem
 - the inconsistent analysis problem
 - the cascading rollback problem
 - deadlocks
 - timeouts
 - deadlock detection