

Advanced Databases

Distributed Transactions

Dr. George Mertzios
Michaelmas Term

george.mertzios@durham.ac.uk

Room 2066, MCS Building

Tel: 42 429

Course Outline


- Enhanced Entity-Relationship (EER) Model
- Semistructured Databases - XML
- XML Data Manipulation - XPath, XQuery
- Transactions and Concurrency Control
- **Distributed Transactions**
- Distributed Concurrency Control

Distributed databases (recap)

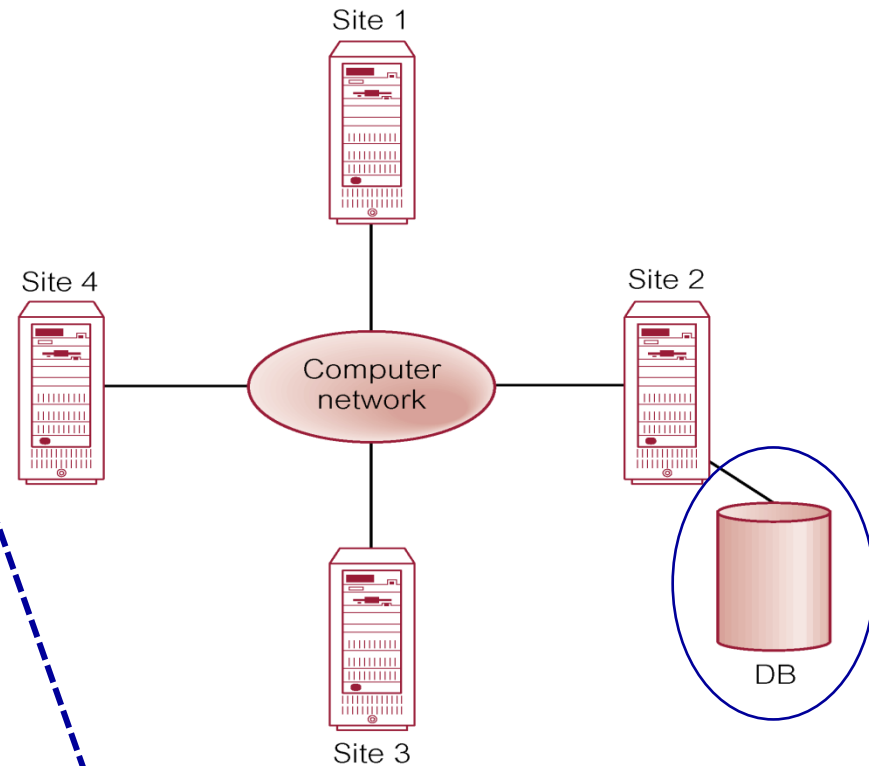
- **Distributed database:**
 - a collection of **shared data**, **distributed** over a **network**
- **Distributed-DBMS (DDBMS):**
 - the **software system** managing the distributed database
- **In a DDBMS:**
 - a **single logical database**, **split** into **fragments**
 - each **fragment** is stored on one (or more) **sites**
 - sites have **local autonomy** (using their own DBMS)
 - sites have access to the **global database**
(using their network connections to other sites)

Distributed databases (recap)

distributed DBMS:



distributed processing
(no distributed DBMS):



Distributed databases (recap)

- A distributed database can be:
 - partitioned
 - database partitioned into disjoint fragments
 - each data item assigned to **exactly one site** (no replication!)
 - no data redundancy
 - completely replicated
 - complete copy of the database at each site
 - allows faster data retrieval
 - selectively replicated
 - combination of partitioning and replication

Distributed Transactions

- A **transaction** in a **DDBMS**:
 - is **initiated** in one of the sites
 - is divided into **sub-transactions** (one for each site)
- The DDBMS must ensure:
 - **synchronization** of **sub-transactions** with other **local transactions**
 - **ACID properties** of (local / global) transactions
- All “**centralized**” **problems** still exist:
 - lost update, uncommitted dependency, inconsistent analysis
- But also a new one appears:
 - **multiple-copy consistency problem**
 - when an item is **updated**, it must be updated at **every site**
 - otherwise inconsistent global database

Distributed Serializability

- The notions of **schedule** and **serializability**:
 - naturally extend to the distributed environment
 - **local schedule** (of sub-transactions) vs. **global schedule**
- A **global schedule** is **serializable** if:
 - each local schedule is serializable (at each site), and
 - the **local serialization orders** (of transactions) are **identical**

In other words:

- all **sub-transactions** appear at **every site**
in the **same order** in the equivalent **serial schedule**

Distributed Serializability

That is:

- n sites S_1, S_2, \dots, S_n
- Denote by T_i^x the sub-transaction of T_i at site S_x
- A **global schedule** of transactions is **serializable** if:
 - whenever $T_i^x < T_j^x$ at some **site** S_x ,
 - we have that $T_i^y < T_j^y$ for **every site** S_y

In other words:

- all **sub-transactions** appear at **every site**
in the **same order** in the equivalent **serial schedule**

Distributed Serializability

Consider a quite restrictive schedule of transactions:

- every site ensures **local serializability**
 - i.e. there exists an equivalent local serial schedule at that site
- at every time at **most one** transaction is **active**

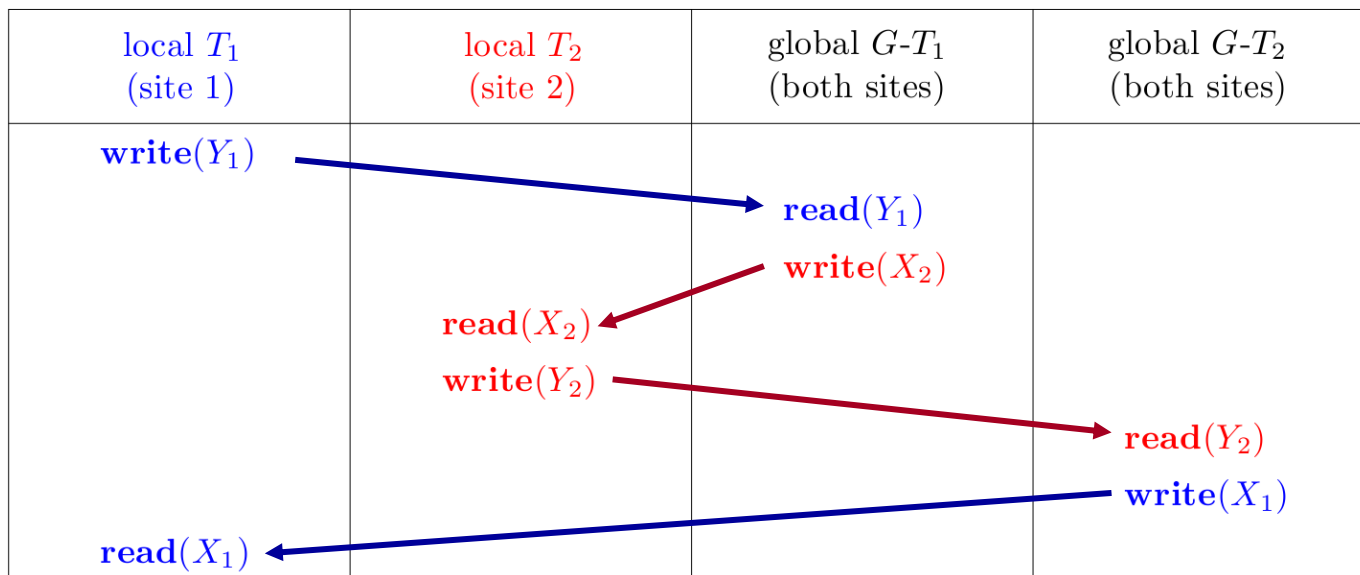
Even then:

- global serializability is **not guaranteed !**

local T_1 (site 1)	local T_2 (site 2)	global $G-T_1$ (both sites)	global $G-T_2$ (both sites)
write (Y_1)		read (Y_1) write (X_2)	
	read (X_2) write (Y_2)		read (Y_2) write (X_1)
read (X_1)			

Distributed Serializability

- Local serializability at **site 1**:
 - $G-T_2 \rightarrow T_1 \rightarrow G-T_1$ is (unique) serialized order
 - Local serializability at **site 2**:
 - $G-T_1 \rightarrow T_2 \rightarrow G-T_2$ is (unique) serialized order
 - Thus, globally:
 - $G-T_1 \rightarrow T_2 \rightarrow G-T_2 \rightarrow T_1 \rightarrow G-T_1$ is not a serialized order
- \Rightarrow the **global schedule** is **not serializable**



Distributed Serializability

Another restrictive schedule of transactions:

- every site ensures **local serializability**
- every **global** transaction is **read-only**

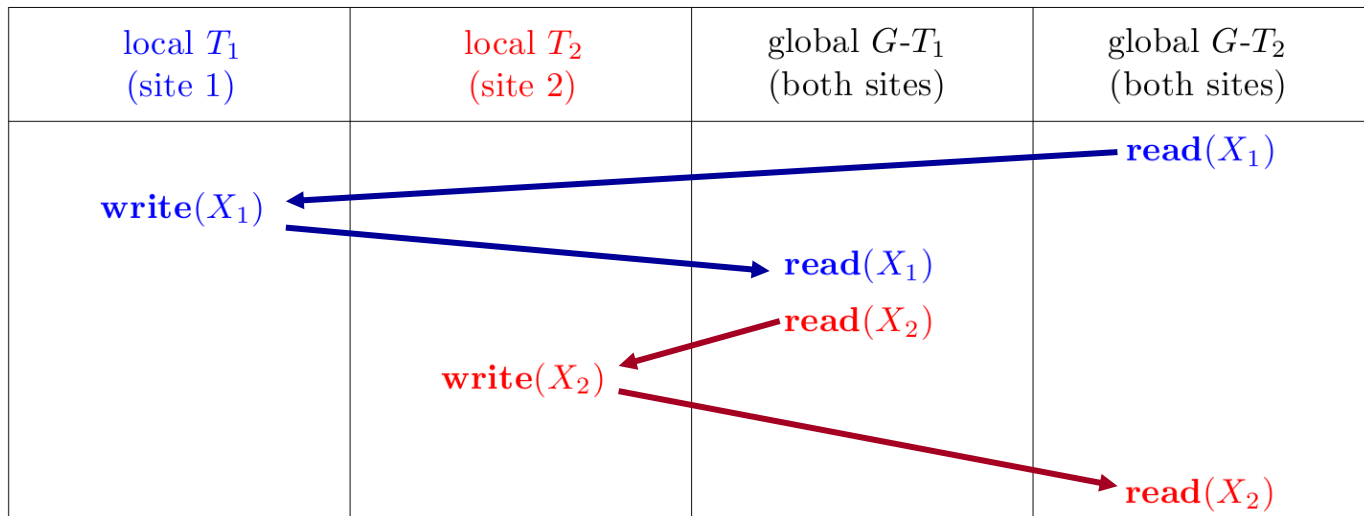
Even then:

- **global serializability** is **not guaranteed** !

local T_1 (site 1)	local T_2 (site 2)	global $G-T_1$ (both sites)	global $G-T_2$ (both sites)
write (X_1)	write (X_2)	read (X_1) read (X_2)	read (X_1) read (X_2)

Distributed Serializability

- Local serializability at **site 1**:
 - $G-T_2 \rightarrow T_1 \rightarrow G-T_1$ is (unique) serialized order
 - Local serializability at **site 2**:
 - $G-T_1 \rightarrow T_2 \rightarrow G-T_2$ is (unique) serialized order
 - Thus, globally:
 - $G-T_1 \rightarrow T_2 \rightarrow G-T_2 \rightarrow T_1 \rightarrow G-T_1$ is not a serialized order
- \Rightarrow the **global schedule** is **not serializable**



Distributed Serializability

Given a distributed non-serial schedule:

- we can (in principle) test conflict serializability using the precedence graph
- the schedule is serializable if and only if the precedence graph has no directed cycle
 - i.e. same as in centralized schedules

What is the technical problem?

- database is distributed


⇒ no site has full information about all (global) conflicts

⇒ even building the precedence graph is not trivial


Distributed Serializability

Simple example:

- the **local** schedule at every site is **serializable**
 - no directed cycles in local precedence graphs
- but the **global** schedule is **not serializable**
 - directed cycle in global precedence graph




site S_1



site S_2

Local
schedules



Global schedule

Distributed Concurrency Control

- How can we **ensure serializability** in practice?
- If all transactions are **local** and **no replication**:
 - all centralized (local) concurrency protocols work
 - **two phase locking (2PL)**; **growing / shrinking phase**
 - **timestamping**
 - otherwise we need to extend them
- Distributed protocols for 2PL:
 - **Centralized 2PL**
 - **Primary Copy 2PL**
 - **Distributed 2PL**
 - **Majority 2PL**

Centralized 2PL

- Centralized 2PL protocol:
 - a single site has a central lock manager (LM)
 - LM maintains all locking information for the DDBMS
- When a transaction is initiated at site S_i :
 - the local transaction coordinator (TC) at S_i is responsible for ensuring consistency throughout the transaction
 - ensure that all copies of an updated item are synchronized
 - if the transaction needs to update a data item x :
 - TC requests from LM a write-lock for every copy of x
 - LM decides to grant the lock or not (standard 2PL rules)
 - similarly for reading a data item x :
 - transaction can read from any copy of x
 - LM decides to grant read-lock or not

Centralized 2PL

- Main idea of centralized 2PL:
 - treat the database as if it were centralized
- Advantages:
 - implementation is easy
 - practically no distributed considerations
 - deadlock detection is simple:
 - build the wait-for graph of DDBMS at central LM
- Disadvantages:
 - bottlenecks
 - overloaded central LM – scalability issues
 - low reliability
 - failure of central LM freezes the whole DDBMS

Primary Copy 2PL

- **Primary Copy 2PL protocol:**
 - straightforward extension of Centralized 2PL
 - **many lock managers (LMs)** across the DDBMS
 - each LM responsible for locking a **different set** of **data items**
- For every **replicated data item x** :
 - one copy is chosen as **primary copy**
 - further copies are **slave copies**
- When an **item x** is **updated**:
 - local TC locates the **primary copy** of x
 - then sends **write-lock request** to the appropriate **LM**
 - only the **primary copy** of x needs to be **locked / updated**
 - later the change **propagates** to **slave copies**

Primary Copy 2PL

- Protocol is very efficient when:
 - large updates are infrequent (high locality of reference)
 - sites do not always need a most updated version of data
- Bottleneck problems are solved:
 - load is distributed to many Lock Managers (LMs)
- Reliability problems still remain:
 - large degree of centralization
 - failure of one LM freezes some part of DDBMS
 - all primary copies of this LM are inaccessible
- Solution:
 - each LM nominates a backup site

Primary Copy 2PL – Backup site

- When an LM receives an update request:
 - LM sends a copy of request to its backup site (B-LM)
- if LM does not send a quick update notification, then the B-LM:
 - assumes that LM failed \Rightarrow acts in place of LM
 - sends a copy of request to its own backup (!)
 - notifies everybody that it is the new LM
 - performs all updates of the original LM
- when LM recovers:
 - it notifies everybody that it is again the LM
 - it receives from B-LM the log of updates made

Distributed 2PL

- Distributed 2PL protocol:
 - one lock manager (LM) at every site
 - managing the locks for the data at this site
- If data is not replicated: same as Primary Copy 2PL
- Otherwise Read-One-Write-All (ROWA) rule:
 - only a read-lock at one site that keeps the item
 - a write-lock at every site that keeps the item
 - the copies in these sites must be locked before the update
- How to check whether a write-lock can be granted?
 - not trivial: high communication cost needed
 - requesting site waits for confirmation from all sites that keep the item

Distributed 2PL

Distributed 2PL

- high communication costs
- always current values

Primary Copy 2PL

- low communication costs
- potentially outdated values
- when needed, every value can synchronize with primary

- Majority 2PL:

- special case of Distributed 2PL

- write-lock is granted if at least half of sites confirm it

⇒ lock holder notifies everybody that it has the lock

- A read-lock: can be simultaneously held by many users
- A write-lock: can be held by only one user each time

Why?

Summary of the Lecture

- Distributed Databases and DDBMSs
- Distributed transactions
- Distributed serializability
- Distributed 2PL concurrency control protocols
 - Centralized 2PL
 - Primary Copy 2PL
 - Distributed 2PL
 - Majority 2PL